

DTCC

2014中国数据库技术大会

DATABASE TECHNOLOGY CONFERENCE CHINA 2014



大数据技术探索和价值发现

Spark Runtime Internals

连城

lian@databricks.com

cheng.lian@ciilabs.org



What is *Spark*

- A fast and general engine for large-scale data processing
- An open source implementation of Resilient Distributed Datasets (RDD)
- Has an advanced DAG execution engine that supports cyclic data flow and *in-memory* computing

Why Spark

- Fast
 - Run machine learning like iterative programs up to 100x faster than Hadoop MapReduce in memory, or 10x faster on disk
 - Run HiveQL compatible queries 100x faster than Hive (with Shark/Spark SQL)

Why Spark

- Easy to use
 - Fluent Scala/Java/Python API
 - Interactive shell
 - 2-5x less code (than Hadoop MapReduce)

Why Spark

- Easy to use
 - Fluent Scala/Java/Python API
 - Interactive shell
 - 2-5x less code (than Hadoop MapReduce)

```
sc.textFile("hdfs://...")  
  .flatMap(_.split(" "))  
  .map(_ -> 1)  
  .reduceByKey(_ + _)  
  .collectAsMap()
```

Why Spark

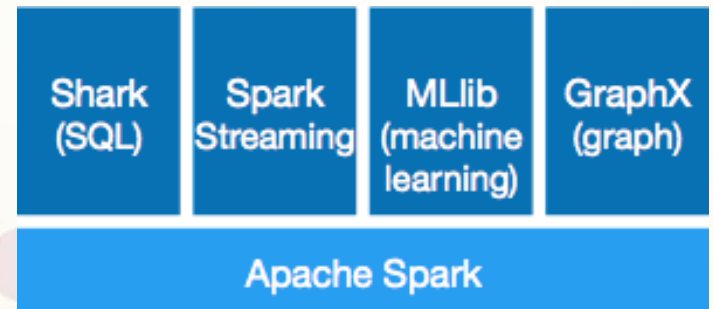
- Easy to use
 - Fluent Scala/Java/Python API
 - Interactive shell
 - 2-5x less code (than Hadoop MapReduce)

```
sc.textFile("hdfs://...")  
  .flatMap(_.split(" "))  
  .map(_ -> 1)  
  .reduceByKey(_ + _)  
  .collectAsMap()
```

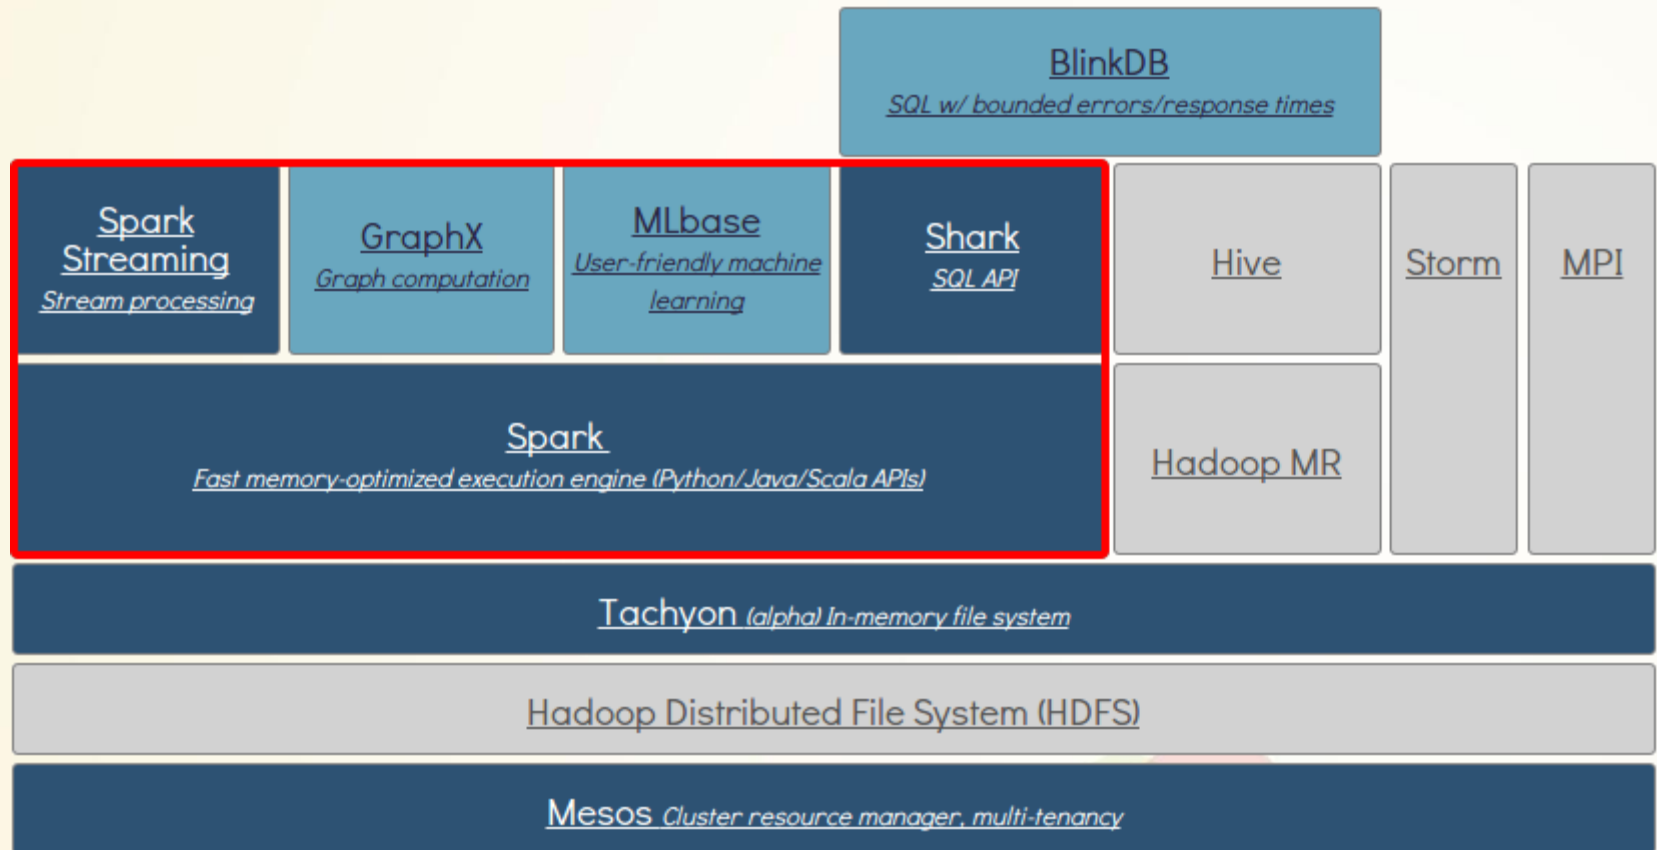
Can you write down
WordCount
in 30 seconds with
Hadoop MapReduce?

Why *Spark*

- Unified big data pipeline for:
 - Batch/Interactive (Spark Core vs MR/Tez)
 - SQL (Shark/Spark SQL vs Hive)
 - Streaming (Spark Streaming vs Storm)
 - Machine learning (MLlib vs Mahout)
 - Graph (GraphX vs Giraph)



A Bigger Picture...



Supported Release

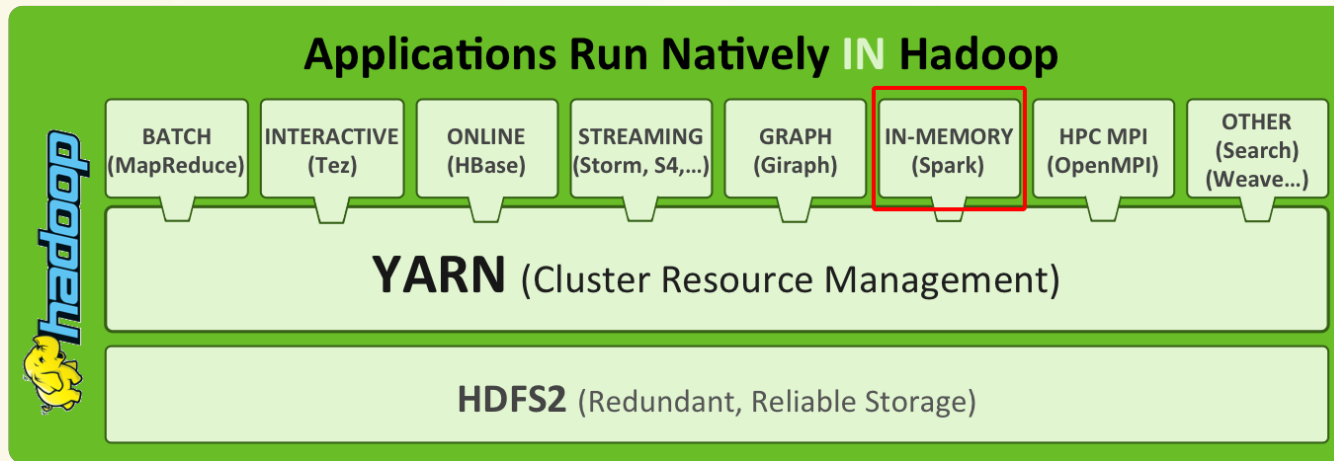


In Development



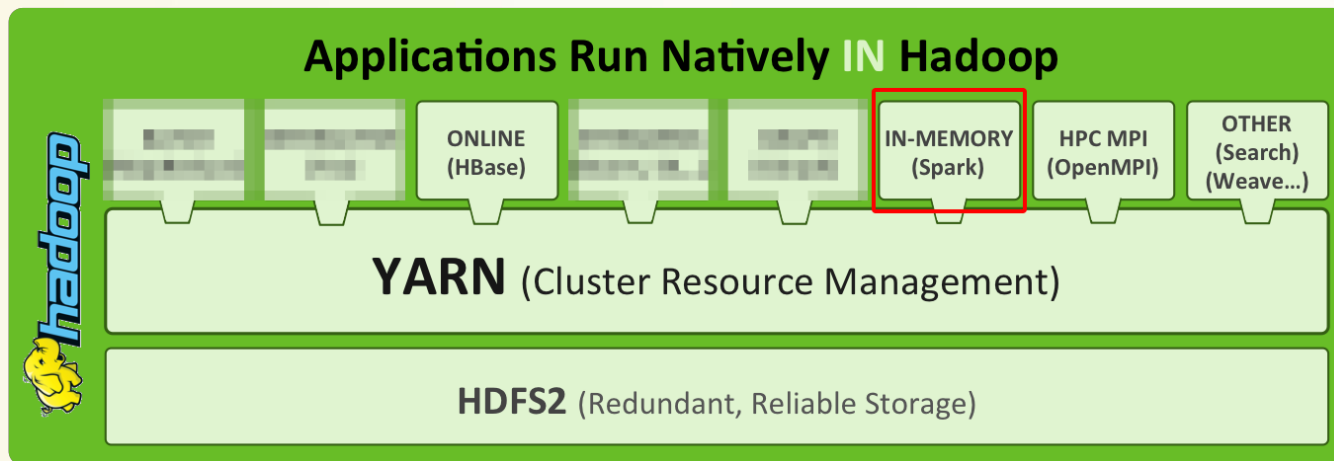
Related External Project

An Even Bigger Picture...




- Components of the Spark stack focus on big data analysis and are compatible with existing Hadoop storage systems
- Users don't need to suffer expensive ETL cost to use the Spark stack

“One Stack To Rule Them All”



- Well, mostly :-)
- And don't forget Shark/Spark SQL vs Hive

Resilient Distributed Datasets



Resilient Distributed Datasets

- Conceptually, RDDs can be roughly viewed as partitioned, locality aware distributed vectors
- An RDD...
 - either points to a direct data source
 - or applies some transformation to its parent RDD(s) to generate new data elements
 - Computation can be represented by *lazy evaluated* lineage *DAGs* composed by connected RDDs

Resilient Distributed Datasets

- Frequently used RDDs can be materialized and cached in-memory to accelerate computation
- Spark scheduler takes data locality into account

The In-Memory Magic

- “In fact, one study [1] analyzed the access patterns in the Hive warehouses at Facebook and discovered that for the vast majority (96%) of jobs, the entire inputs could fit into a fraction of the cluster’s total memory.”

[1] G. Ananthanarayanan, A. Ghodsi, S. Shenker, and I. Stoica. Disk-locality in datacenter computing considered irrelevant. In HotOS '11, 2011.

The In-Memory Magic


- Without cache
 - Elements are accessed in an iterator-based streaming style
 - One element a time, no bulk copy
 - Space complexity is almost $O(1)$ when there's only narrow dependencies

The In-Memory Magic


- With cache
 - One block per RDD partition
 - LRU cache eviction
 - Locality aware
 - Evicted blocks can be *recomputed in parallel* with the help of RDD lineage DAG

Play with Error Logs

```
sc.textFile("hdfs://<input>")  
  .filter(_.startsWith("ERROR"))  
  .map(_.split(" ")(1))  
  .saveAsTextFile("hdfs://<output>")
```




Step by Step




The `RDD.saveAsTextFile()` action triggers a job. Tasks are started on scheduled executors.

Step by Step



The task pulls data lazily from the final RDD (the MappedRDD here).

Step by Step



A record (text line) is pulled out from HDFS...

Step by Step

INFO ...



... into a HadoopRDD

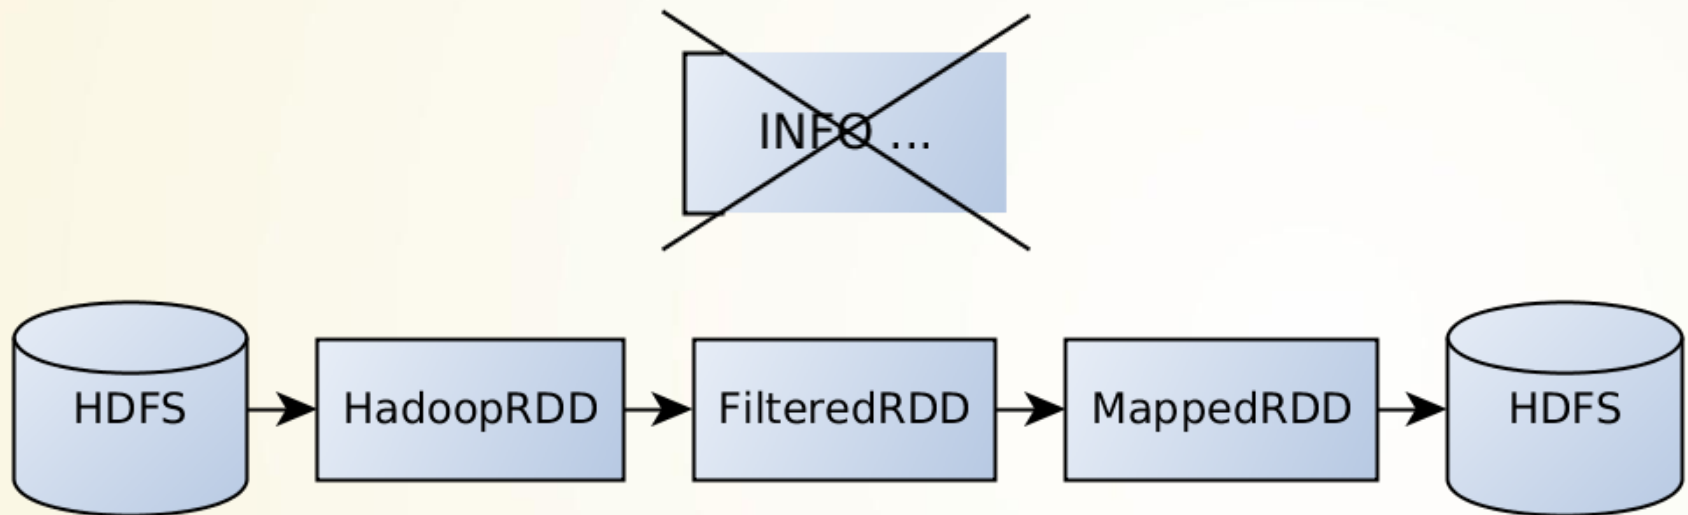
Step by Step

INFO ...



Then filtered with the FilteredRDD

Step by Step



Oops, not a match

Step by Step

ERROR ...



Another record is pulled out...

Step by Step

ERROR ...



Filtered again...

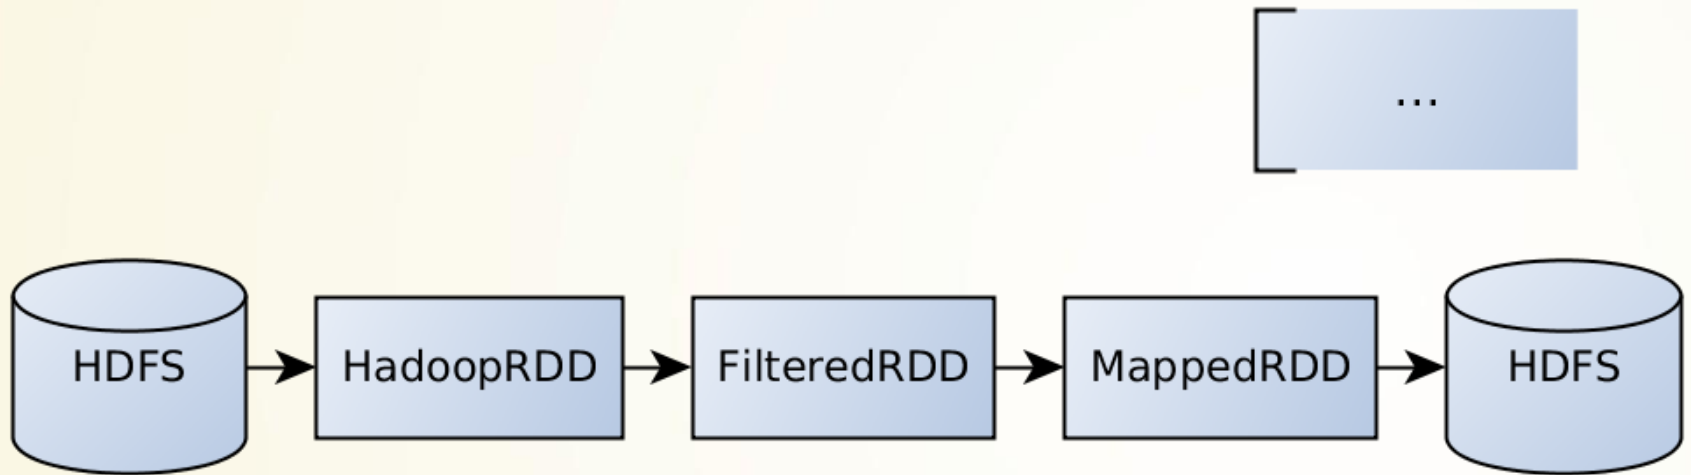
Step by Step

ERROR ...




Passed to the MappedRDD for next transformation

Step by Step



Transformed by the MappedRDD (the error message is extracted)

Step by Step



Saves the message to HDFS

A Synthesized Iterator

```
new Iterator[String] {  
  private var head: String = _  
  private var headDefined: Boolean = false  
  
  def hasNext: Boolean = headDefined || {  
    do {  
      try head = readOneLineFromHDFS(...) // (1) read from HDFS  
      catch {  
        case _: EOFException => return false  
      }  
    } while (!head.startsWith("ERROR")) // (2) filter closure  
    true  
  }  
  
  def next: String = if (hasNext) {  
    headDefined = false  
    head.split(" ")(1) // (3) map closure  
  } else {  
    throw new NoSuchElementException("...")  
  }  
}
```

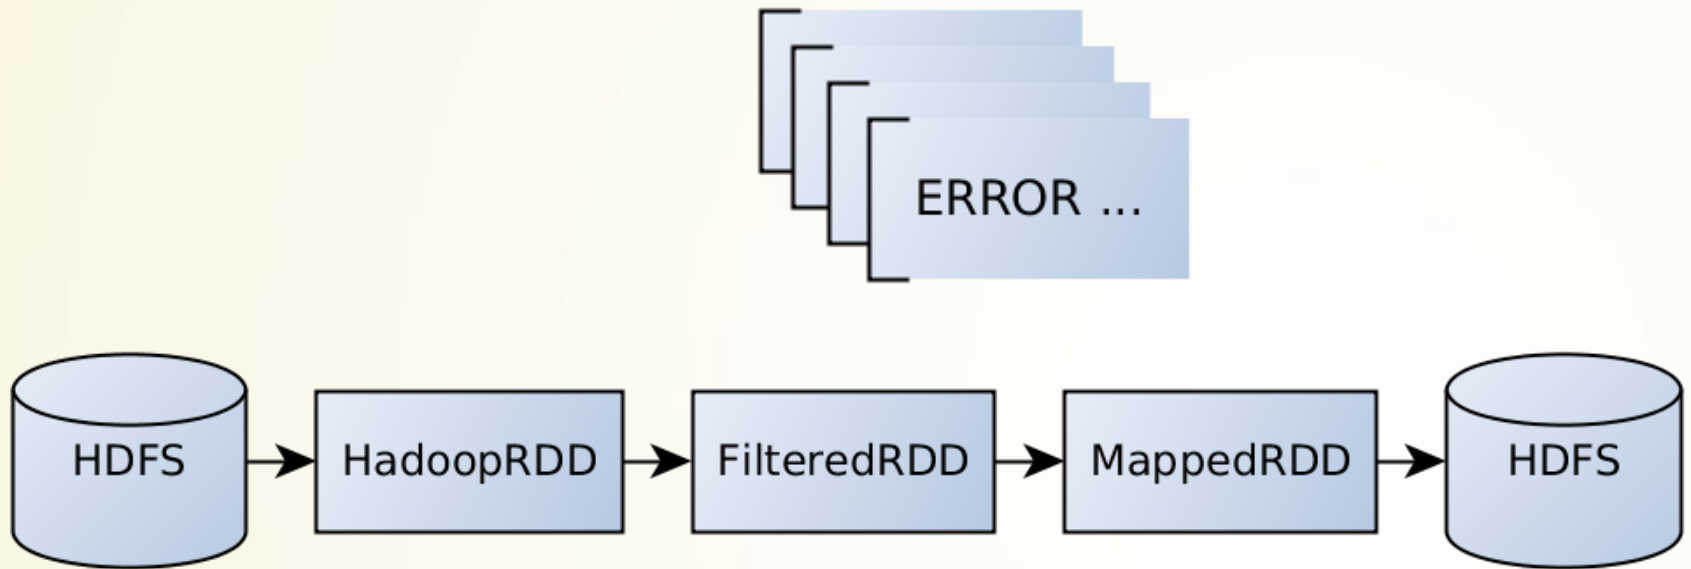
Constant Space
Complexity!

If Cached...

```
val cached = sc
  .textFile("hdfs://<input>")
  .filter(_.startsWith("ERROR"))
  .cache()

cached
  .map(_._split(" ")(1))
  .saveAsTextFile("hdfs://<output>")
```

If Cached...




All filtered error messages are cached in memory before being passed to the next RDD. LRU cache eviction is applied when memory is insufficient

But... I Don't Have Enough
Memory To Cache *All* Data


Don't Worry

- In most cases, you only need to cache a fraction of hot data extracted/transformed from the whole input dataset
- Spark performance downgrades linearly & gracefully when memory decreases



How Does The Job Get Distributed?

Spark Cluster Overview



Driver, Cluster Manager,
Worker, Executor, ...

So... Where THE HELL
Does My Code Run?

Well... It Depends

Driver and SparkContext

```
import org.apache.spark.SparkContext  
import org.apache.spark.SparkContext._
```

```
val sc = new SparkContext(...)
```

- A **SparkContext** initializes the application driver, the latter then registers the application to the cluster manager, and gets a list of executors
- Since then, the driver takes full responsibilities

WordCount Revisited

```
val lines = sc.textFile("input")
val words = lines.flatMap(_.split(" "))
val ones = words.map(_ -> 1)
val counts = ones.reduceByKey(_ + _)
val result = counts.collectAsMap()
```

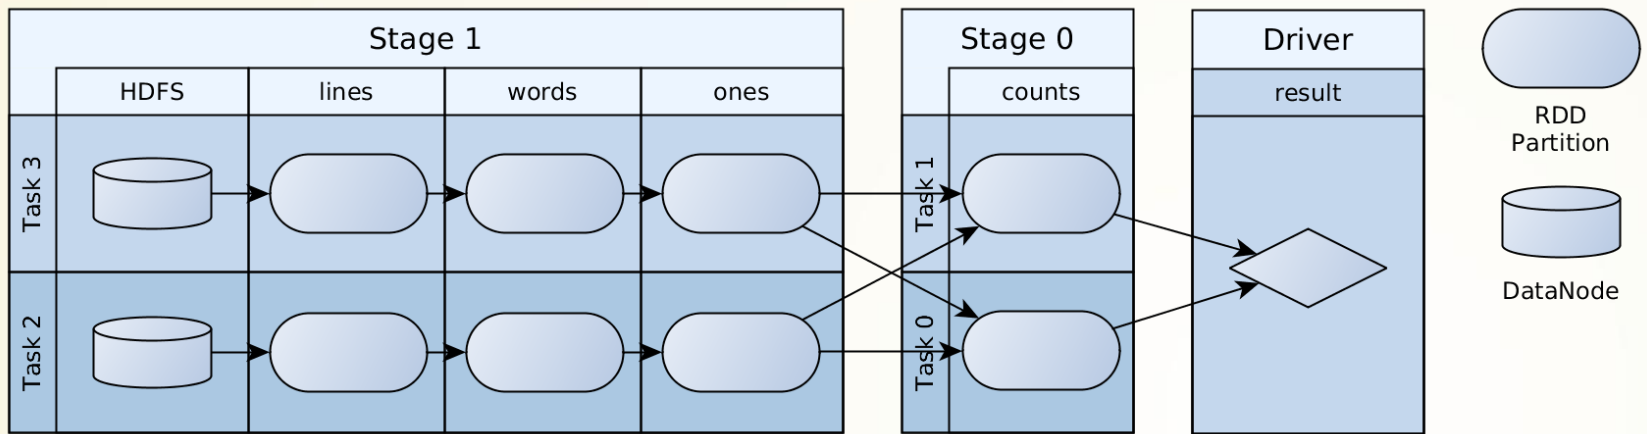
- RDD lineage DAG is built on *driver* side with:
 - Data source RDD(s)
 - Transformation RDD(s), which are created by transformations

WordCount Revisited

```
val lines = sc.textFile("input")
val words = lines.flatMap(_.split(" "))
val ones = words.map(_ -> 1)
val counts = ones.reduceByKey(_ + _)
val result = counts.collectAsMap()
```


- Once an **action** is triggered on *driver* side, a job is submitted to the *DAG scheduler* of the driver

WordCount Revisited



- DAG scheduler cuts the DAG into stages and turns each *partition* of a stage into a single task.
- DAG scheduler decides *what to run*

WordCount Revisited



- Tasks are then scheduled to executors by driver side *task scheduler* according to *resource* and *locality* constraints
- Task scheduler decides *where to run*

WordCount Revisited

```
val lines = sc.textFile("input")
val words = lines.flatMap(_.split(" "))
val ones = words.map(_ -> 1)
val counts = ones.reduceByKey(_ + _)
val result = counts.collectAsMap()
```

- Within a task, the **lineage DAG** of **corresponding stage** is serialized together with **closures** of transformations, then sent to and executed on scheduled *executors*

WordCount Revisited

```
val lines = sc.textFile("input")
val words = lines.flatMap(_.split(" "))
val ones = words.map(_ -> 1)
val counts = ones.reduceByKey(_ + _)
val result = counts.collectAsMap()
```

- The **reduceByKey** transformation introduces in a shuffle
- Shuffle outputs are written to local FS on the mapper side, then downloaded by reducers

WordCount Revisited

```
val lines = sc.textFile("input")  
val words = lines.flatMap(_.split(" "))  
val ones = words.map(_ -> 1)  
val counts = ones.reduceByKey(_ + _)  
val result = counts.collectAsMap()
```

- **ReduceByKey** automatically combines values within a single partition locally on the mapper side and then reduce them globally on the reducer side.

WordCount Revisited

```
val lines = sc.textFile("input")
val words = lines.flatMap(_.split(" "))
val ones = words.map(_ -> 1)
val counts = ones.reduceByKey(_ + _)
val result = counts.collectAsMap()
```

- At last, results of the action are sent back to the driver, then the job finishes.

What About Parallelism?

How To Control Parallelism?

- Can be specified in a number of ways
 - RDD partition number
 - `sc.textFile("input", minSplits = 10)`
 - `sc.parallelize(1 to 10000, numSlices = 10)`
 - Mapper side parallelism
 - Usually inherited from parent RDD(s)
 - Reducer side parallelism
 - `rdd.reduceByKey(_ + _, numPartitions = 10)`
 - `rdd.reduceByKey(partitioner = p, _ + _)`
 - ...

How To Control Parallelism?

- “Zoom in/out”
 - `RDD.repartition(numPartitions: Int)`
 - `RDD.coalesce(numPartitions: Int, shuffle: Boolean)`


A Trap of Partitions

```
sc.textFile("input", minSplits = 2)
  .map { line =>
    val Array(key, value) = line.split(",")
    key.toInt -> value.toInt
  }
  .reduceByKey(_ + _)
  .saveAsText("output")
```

- Run in *local* mode
- 60+GB input file

A Trap of Partitions

```
sc.textFile("input", minSplits = 2)
  .map { line =>
    val Array(key, value) = line.split(",")
    key.toInt -> value.toInt
  }
  .reduceByKey(_ + _)
  .saveAsText("output")
```



A Trap of Partitions

```
sc.textFile("input", minSplits = 2)
  .map { line =>
    val Array(key, value) = line.split(",")
    key.toInt -> value.toInt
  }
  .reduceByKey(_ + _)
  .saveAsText("output")
```

- Actually, this may generate **4,000,000** temporary files...
- Why 4M? (Hint: $2K^2$)



A Trap of Partitions

```
sc.textFile("input", minSplits = 2)
  .map { line =>
    val Array(key, value) = line.split(",")
    key.toInt -> value.toInt
  }
  .reduceByKey(_ + _)
  .saveAsText("output")
```

- 1 partition per HDFS split
- Similar to mapper task number in Hadoop MapReduce, the **minSplits** parameter is only a *hint* for split number

A Trap of Partitions

```
sc.textFile("input", minSplits = 2)
  .map { line =>
    val Array(key, value) = line.split(",")
    key.toInt -> value.toInt
  }
  .reduceByKey(_ + _)
  .saveAsText("output")
```

- Actual split number is also controlled by some other properties like minSplitSize and default FS block size

A Trap of Partitions

```
sc.textFile("input", minSplits = 2)
  .map { line =>
    val Array(key, value) = line.split(",")
    key.toInt -> value.toInt
  }
  .reduceByKey(_ + _)
  .saveAsText("output")
```

- In this case, the final split size equals to local FS block size, which is 32MB by default, and $60\text{GB} / 32\text{MB} \approx 2\text{K}$
- ReduceByKey generates 2K^2 shuffle outputs

A Trap of Partitions

```
sc.textFile("input", minSplits = 2).coalesce(2)
  .map { line =>
    val Array(key, value) = line.split(",")
    key.toInt -> value.toInt
  }
  .reduceByKey(_ + _)
  .saveAsText("output")
```

- Use **RDD.coalesce()** to control partition number precisely.

Summary

- In-memory magic
 - Iterator-based streaming style data access
 - Usually you don't need to cache all dataset
 - When memory is not enough, Spark performance downgrades linearly & gracefully

Summary

- Distributed job
 - Locality aware
 - Resource aware
- Parallelism
 - Reasonable default behavior
 - Can be finely tuned

References

- [Resilient distributed datasets: a fault-tolerant abstraction for in-memory cluster computing](#)
- [An Architecture for Fast and General Data Processing on Large Clusters](#)
- Spark Internals ([video](#), [slides](#))
- [Shark: SQL and Rich Analytics at Scale](#)
- [Spark Summit 2013](#)

Q&A

THANKS

